

Azure Storage

August 2021

Camilla Gaardsted, SuperUsers A/S

Data



Raspberry Pi 4 Model B 4GB DDR4 RAM

Varenummer: **144-155**

Varekode: RPI4-MODBP-4GB

Lagervare!



Tesla, Inc.
NASDAQ: TSLA

705,67 USD +10,89 (1,57 %) ↑

Lukket: 4. jan. 07.14 GMT-5 · Ansvarsfraskrivelse
Uofficiel handel 722,50 +16,83 (2,38 %)

1 dag 5 dage 1 måned 6 måneder ÅTD



Name	Date modified	Type	Size
ERRORLOG	04-01-2021 02:12	File	56.920 KB
ERRORLOG.1	09-12-2020 07:00	1 File	59.995 KB
ERRORLOG.2	11-11-2020 07:03	2 File	43.463 KB
ERRORLOG.3	19-10-2020 12:53	3 File	65 KB
ERRORLOG.4	19-10-2020 11:05	4 File	38 KB
ERRORLOG.5	19-10-2020 11:01	5 File	38 KB
ERRORLOG.6	19-10-2020 10:55	6 File	238 KB



Donald J. Trump
@realDonaldTrump

Following

The concept of global warming was created by and for the Chinese in order to make U.S. manufacturing non-competitive.

RETWEETS
104,728

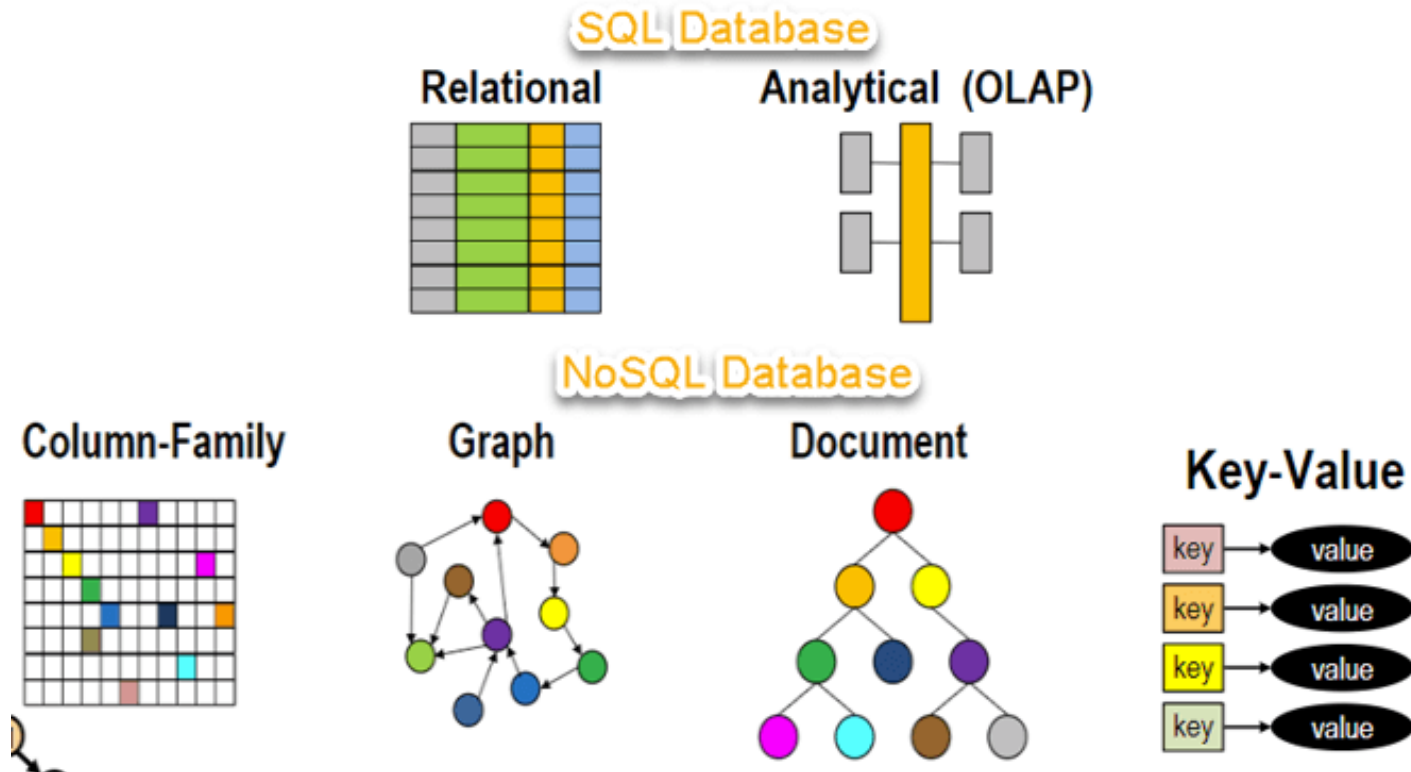
LIKES
67,204



7:15 PM - 6 Nov 2012

12K 105K 67K

NoSQL Nye måder at gemme data på




Data filformater

- CSV
- APACHE PARQUET
- AVRO
- JSON

Dataset	Size on Amazon S3	Query Run Time	Data Scanned	Cost
Data stored as CSV files	1 TB	236 seconds	1.15 TB	\$5.75
Data stored in Apache Parquet Format	130 GB	6.78 seconds	2.51 GB	\$0.01
Savings	87% less when using Parquet	34x faster	99% less data scanned	99.7% savings

Kilde:<https://databricks.com/glossary/what-is-parquet>

Classic ETL




Big Data

Big data er et begreb indenfor datalogi, der bredt dækker over indsamling, opbevaring, analyse, processering og fortolkning af enorme mængder af data [Kilde: https://da.wikipedia.org/wiki/Big_data]



Modern ELT



Hadoop, Spark and HDInsights

Parallel distributed processing of data


- Hadoop
- Spark Cluster
- Kafka
- Azure HDInsights



SQL Server Integration Services

SQL Server Integration Services (SSIS, 2005+)

- ETL Tool from MS distributed with MS SQL Server
- Jobs/packages for data movement, transformations, processing, backup/restore
- Run once/scheduled/ad hoc



Azure Data Factory (ADF)

- Cloud udgaven/erstatningen af SSIS
- Pipelines i ADF svarer til pakker i SSIS
- Kan køre SSIS pakker

AI og ML

Artificial Intelligence (AI)

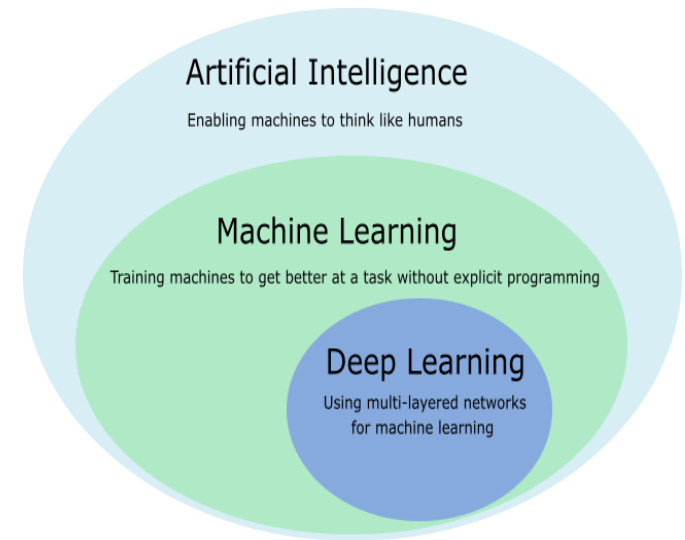
- Perform tasks normally requiring human intelligence

Machine Learning (ML)

- Computer science + **statistics**
- **Feed** the algorithm with data
- Recognize patterns in data
- Predictions for new data


Deep Learning

- Artificial neural network (hidden layers – deeper)
- Algorithms analyzes data using a logic structure like humans




Power BI

Portefølje af produkter og services til at transformere, visualisere og præsentere data i **interaktive** rapporter.



Azure

Store, transform, process, analyze, and visualize data



Azure Data storage

Azure ressource navn	Type
Storage account	Blob
Storage account	Fileshare
Storage account	Table storage
Storage account	Queue storage
Storage account	Data Lake
(CosmosDB)	Gremlin, SQL, etc
(Azure SQL Database)	Relationel database
Azure Synapse Analytics	Data warehouse + mere
Azure Event Hub	
(Azure IoT Hub)	

Azure data processing

Teknologi	Input	Output
Azure Stream Analytics	Event/IoT Hub, Datalake	Blob, Power BI stream
Azure Databricks	Alt (Python til råddighed)	
Azure HDInsights		
Azure Synapse Analytics	Alt (Python til råddighed)	
Azure Data Factory		

Data Lake

En sø af data hvor vi bare hælder alle mulige former for rå data ind som **filer**



TXT, CSV
JSON, XML
PDF,
AVRO,
PARQUET

Weather
Tweets
Images
Logs
Audio
Videos
etc

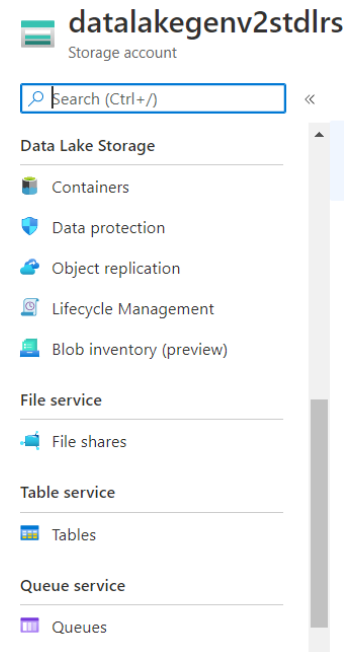
Azure Data Lake

Data Lake Storage Gen1 (**deprecated**, own resource)

Data Lake Storage Gen2 (via STD storage account)

- Enable hierarchical namespace (**on creation**)
- All benefits from blob storage

Blob service section just replaced with
Data Lake Storage section



Azure Data Lake Gen2

Hierarchical namespace

- Rename/delete directory is a metadata operation
- Search – can skip irrelevant paths

Hadoop compatible access

- Manage and access data just as you would with a Hadoop Distributed File System (HDFS)

Avoid many small files (100MB to 1GB is optimal)

Partition data in folders used in common filters of the data (year/month/ , country/customer/ etc)

Blob or Azure Data Lake Gen2?

Blobs are stored as objects in one single flat namespace e.g. the slash is just part of the name

- data/2021/januarysale.csv

Filesystem in a data lake

- Data
 - 2021
 - Januarysale.csv

Cosmos DB

Cosmos DB is a fully managed **NoSQL database**

- Single-digit millisecond response time
- Automatic and instant scalability

A database belongs to *one* Azure Cosmos DB account with a unique DNS name

- <https://<accountname>.documents.azure.com>

API is determined at account level

One account can have many databases (same API)

Azure Cosmos DB



SQL

SQL

{LEAF}

API for MongoDB



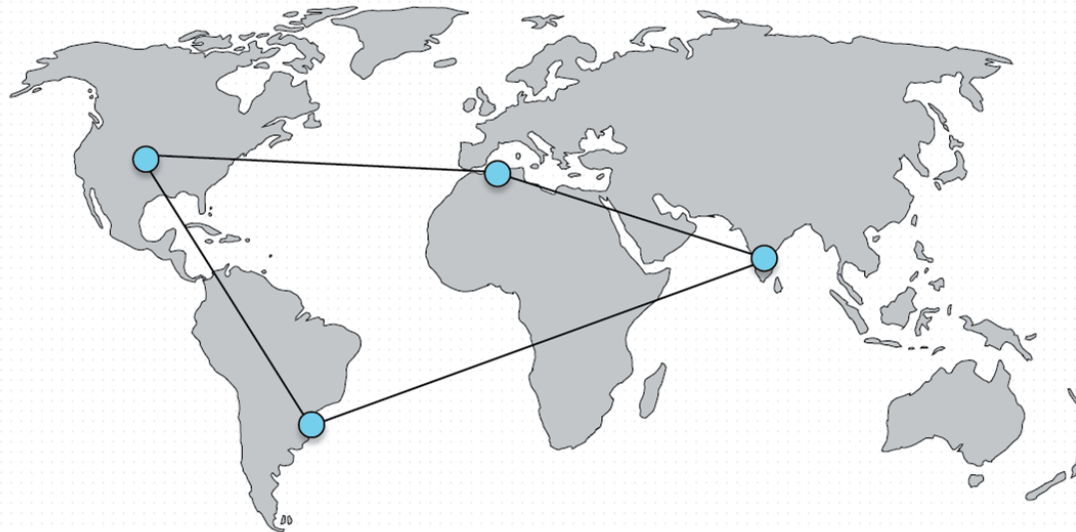
Gremlin



Cassandra



Table




Key-Value



Column-Family



Documents



Graph



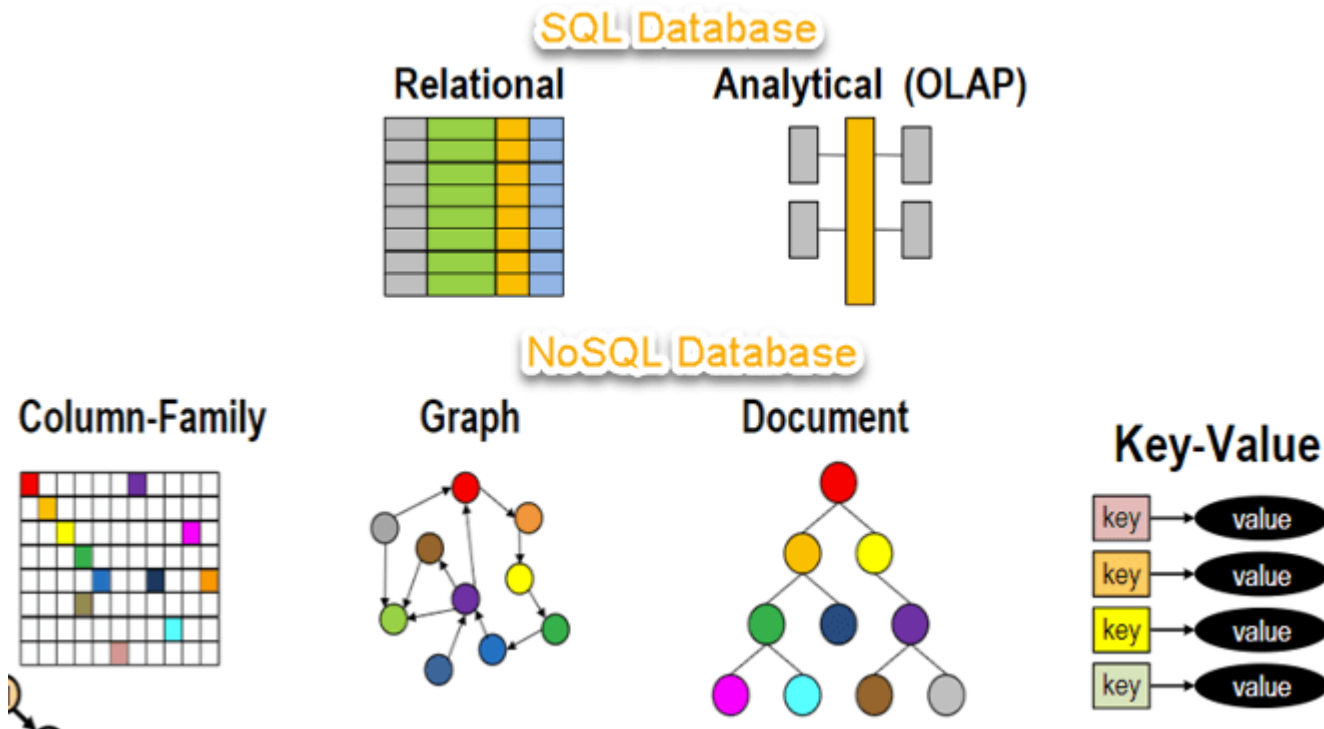
Guaranteed speed at any scale

Simplified application development

Mission-critical ready

Fully managed and cost effective

NoSQL – 4 types of databases



Azure Cosmos DB API

One Azure Cosmos DB Account

5 different APIs (**wire protocol and storage format**)

- **Core (SQL)** (GlobalDocumentDB) (default)
- MongoDB API
- Cassandra
- Azure Table
- Gremlin (graph)


CosmosDB – Capacity mode

Criteria	Provisioned throughput	Serverless
Status	Generally available	In preview
Best suited for	Workloads with sustained traffic requiring predictable performance	Workloads with intermittent or unpredictable traffic and low average-to-peak traffic ratio
How it works	For each of your containers, you provision some amount of throughput expressed in Request Units per second. Every second, this amount of Request Units is available for your database operations. Provisioned throughput can be updated manually or adjusted automatically with autoscale .	You run your database operations against your containers without having to provision any capacity.
Geo-distribution	Available (unlimited number of Azure regions)	Unavailable (serverless accounts can only run in 1 Azure region)
Maximum storage per container	Unlimited	50 GB
Performance	< 10 ms latency for point-reads and writes covered by SLA	< 10 ms latency for point-reads and < 30 ms for writes covered by SLO
Billing model	Billing is done on a per-hour basis for the RU/s provisioned, regardless of how many RUs were consumed.	Billing is done on a per-hour basis for the amount of RUs consumed by your database operations.

Azure Cosmos DB – Backup policy

Defined at Account level

Create Azure Cosmos DB Account

 For a limited time, create a new Azure Cosmos DB account with multi-region writes in any region, and receive up to 33% off for the life of your account. Restrictions apply.*

Basics Networking Backup Policy Encryption Tags Review + create

Azure Cosmos DB provides two different backup policies. You will not be able to switch between backup policies after the account has been created.

Backup policy ⓘ

Periodic Continuous

[Sign up for enabling continuous backup policy](#)

Backup interval ⓘ

60 ✓ Minute(s) ▼
60-1440

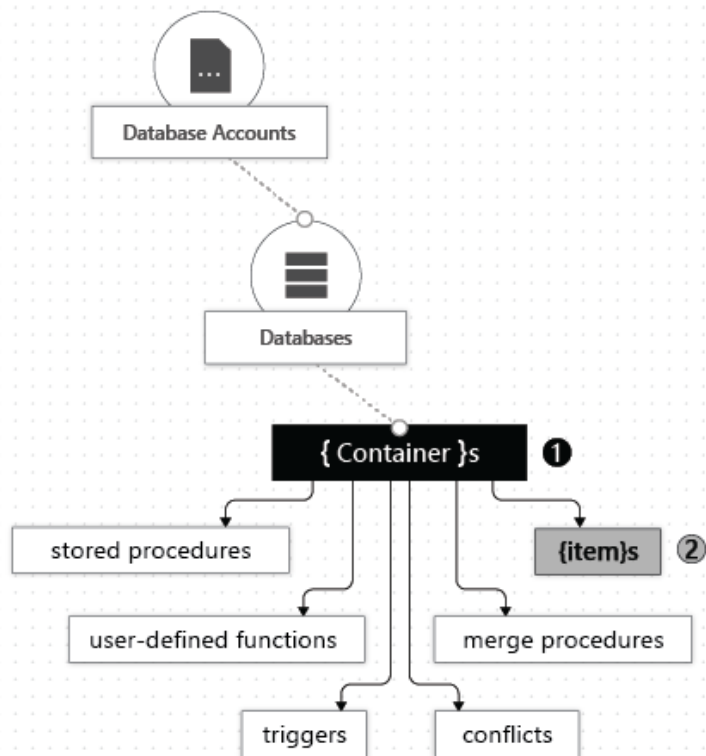
Backup retention ⓘ

8 Hours(s) ▼
8-720

Copies of data retained 8


For additional pricing details, please check [here](#)

CosmosDB – container/items



① { Container }s

Depending on the Cosmos API, a container is realized as:



② { item}s

Depending on the Cosmos API, an item is realized as:



CosmosDB - Container

- A Cosmos container is a schema-free collection of **JSON items**
- An Azure Cosmos container is the unit of scalability both for provisioned throughput and storage.
- A container is horizontally partitioned and then replicated across multiple regions.
- The items that you add to the container and the throughput that you provision on it are automatically distributed across a set of logical partitions based on the partition key.

Cosmos DB Container Partitionkey

Logical partitions are formed based on the value of a *partition key* that is associated with each item in a container.

Each item in a container has an *item ID* (unique within a logical partition)

The partition key and the *item ID* creates the item's *index*, which uniquely identifies the item.

Once you select your partition key, it is not possible to change it in-place.

Cosmos DB – Partition key

For all containers, your partition key should:

- Be a property that has a value which does not change. If a property is your partition key, you can't update that property's value.
- Have a high cardinality (wide range of possible values)
- Spread request unit (RU) consumption and data storage evenly across all logical partitions.

CosmosDB – Request Units (RU)

1 RU = cost to read a 1 KB document

5 RU = cost to write a 1 KB document

Minimum per 400 RU

Provision RU throughput at

- Database level (all containers share the amount)
- Container level (dedicated amount)

View Query Stats in Data Explorer

Possible to mix shared/dedicated, but container cannot change mode later

CosmosDB – Throughput

- Database
- Container

Autopilot(preview)

- $0.1 * T_{max} < T < T_{max}$
- 4 levels – max determines max size

CosmosDB - RU

- While you estimate the number of RUs per second to provision, consider the following factors:
- **Item size:** As the size of an item increases, the number of RUs consumed to read or write the item also increases.
- **Item indexing:** By default, each item is automatically indexed. Fewer RUs are consumed if you choose not to index some of your items in a container.
- **Item property count:** Assuming the default indexing is on all properties, the number of RUs consumed to write an item increases as the item property count increases.
- **Indexed properties:** An index policy on each container determines which properties are indexed by default. To reduce the RU consumption for write operations, limit the number of indexed properties.
- **Data consistency:** The strong and bounded staleness consistency levels consume approximately two times more RUs while performing read operations when compared to that of other relaxed consistency levels.

CosmosDB - Index

Auto indexing for a container. Can be turned off

- Consistent
- None

Including and excluding property paths

Can be set in the portal under Scale & Settings

Composite index is allowed

Cosmos DB – SQL API development

Platform options:

- .NET
- Python
- Java
- Node.js
- Xamarin

dotnet CLI cmd tool example in VSCode

```
dotnet new console
```

Default consistency level

Angives på en Azure Cosmos DB account

Der er 5 muligheder Strong -> Eventual

Eksempler med musiknoder som forklarer levels:

<https://docs.microsoft.com/en-us/azure/cosmos-db/consistency-levels>

Cosmos DB – SQL API

Completely different from ANSI SQL

A document is a JSON item

The result of a query is a valid JSON value

SQL API works on **JSON** values, it deals with **tree-shaped** entities instead of rows and columns

NB Case sensitive and beware of number/string

Refer to the tree nodes at any arbitrary depth, like
Node1.Node2.Node3....Node<n>

Point reads (key/value lookup) vs SQL queries

Azure Cosmos DB – Data Explorer

Data Explorer

The screenshot displays the Azure Cosmos DB Data Explorer interface. The top toolbar includes icons for file operations and actions like 'New Item', 'Update', 'Discard', 'Delete', and 'Upload Item'. The left sidebar shows the 'SQL API' section with a tree view containing 'Users' and 'WebCustomers'. Under 'WebCustomers', the 'Items' collection is selected. The main area shows a table of items with columns 'id' and '/...'. The first item has 'id' 1 and '/...' 'yanhe'. The second item has 'id' 2 and '/...' 'nelapin'. A 'Load more' link is visible below the table. To the right of the table, a JSON representation of the first item is shown, including fields like 'id', 'userId', 'lastName', 'firstName', 'email', 'dividend', and 'OrderHistory'.

SQL API

- Users
 - WebCustomers
 - Items
 - Scale & Settings
 - Stored Procedures
 - User Defined Functions
 - Triggers

Items

SELECT * FROM c

Edit Filter

id	/...
1	yanhe
2	nelapin

Load more

```
1 {  
2   "id": "1",  
3   "userId": "yanhe",  
4   "lastName": "Suhlemia",  
5   "firstName": "Yan",  
6   "email": "yanhe@contoso.com",  
7   "dividend": null,  
8   "OrderHistory": [  
9     {  
10      "OrderId": "1000",  
11      "DateShipped": "08/17/2018",  
12      "Total": "52.49"  
13    }  
14  ],  
15 }
```

Azure Cosmos DB – Data Explorer

Data Explorer is a tool for the Cosmos DB SQL API in the Azure Portal

Browse/View/Create/Delete

- Databases
- Containers

Browse/View/Create/Update/Delete

- Items
- Stored Procedures
- User Defined Functions
- Triggers

Azure Cosmos DB – DB SQL API

.NET Querying (JSON) documents via

- LINQ
- SQL

Java, Python etc API

Cosmos DB – Resource tokens

Azure Cosmos DB uses two types of keys to authenticate users and provide access to its data and resources:

- Primary Keys
- Resource tokens

Used for application resources: containers, documents, attachments, stored procedures, triggers, and UDFs

Azure Cosmos DB – Security Access

Two account keys for **administrative** resources:
database accounts, databases, users, and
permissions

Two account keys for Read-only access on account

Read-write Keys Read-only Keys

URI

`https://cosmos20210108.documents.azure.com:443/`

PRIMARY KEY

`5bl2O8w3lu2TrPRfRw8fn2UdDUXy4Ksc0xlEpcREt5YL6epEI2BamVWnEX6b3w5OFn93mdh8AQN3CpoHhT62Q==`

SECONDARY KEY

`awxCN270AB0qDFdl35xQV8R261zg5QHPCglLNbduN12dQdWpl00PmvJNJrl6UD0pJHd0XHTUZAoFqoFosoRBA==`

Cosmos DB – Database Users

A database can contain zero or more users

Permissions on a resource

- All (full permission)
- Read (no write, update or delete)

Running a stored procedure requires All permission on the container

SQL Databases in Azure

MS SQL 4 options in Azure:

- Azure SQL Database
- VM with MS SQL Server 2017+ standard/enterprise
- Managed instance
- Azure Synapse Analytics (Data warehouse)

Alternativer:

- Azure database for MySQL, PostgreSQL, etc
- VM with Oracle, MySQL, etc

Azure Data Factory - history

Data movement, transformations, processing data, ETL tools:

- SQL Server Integration Services (SSIS, 2005+)
- Azure Data Factory v1 (2015+)
- Azure Data Factory v2 (2018+)

Pipelines i ADF svarer til pakker i SSIS

Azure Data Factory (ADF)

Data integration

Data Processing

ETL and ELT (skema for datamodel)

SSIS integration runtime

Administrate ADF via GUI or json files

The serverless integration service does the rest..



Azure Data Factory - Steps

1. Connect to all the required sources of data and processing, such as software-as-a-service (SaaS) services, databases, file shares, and FTP web services.
2. Move the data to a centralized location for subsequent processing.
3. Process, analyze and/or transform data
4. Export/load data into destination


Azure Data Factory – Git repository

Git integration afgør om man har en **Save** eller **Publish** knap

Git via: Azure DevOps or GitHub

Configure only your development data factory with Git integration.

Git repository

Git repository information associated with your data factory. [CI/CD best practices](#) 

 Setting  Disconnect

Repository type	GitHub
GitHub account	camillagaardsted
Repository name	adf
Collaboration branch	master
Publish branch	adf_publish
Root folder	/

Azure Data Factory

- Linked services (connection to the data source)
- Datasets (structure of the data)
- Pipeline
- Activities
- Data flows
- Triggers
- Integration runtimes (IR)

ADF – Linked services

A linked service is connection information e.g. like a connectionstring

View/edit via ADF->Monitor->Linked services

ADF - Dataset


- Data store
- Format

ADF – Copy Activity

Source options (many)

Sink options (limited)

You can use the Copy activity to copy files as-is between two file-based data stores, in which case the data is copied efficiently without any serialization or deserialization.



ADF - Data flow activity

- Mapping data flow
 - Uses Azure Databricks cluster
 - Visual flow
 - Each step in the flow is a transformation
 - Preview data via Debug
 - Handles also inserts, updates, deletes and upserts
- Wrangling data flow (**omdøbt til Power Query**)
 - Uses a managed Spark environment
 - Power Query Online mashup editor (M)
 - Not all M commands are supported!

ADF – Integration Runtime (IR)

IR type

Azure

Public network

Data Flow

Data movement

Activity dispatch

Private network

Self-hosted

Data movement

Activity dispatch

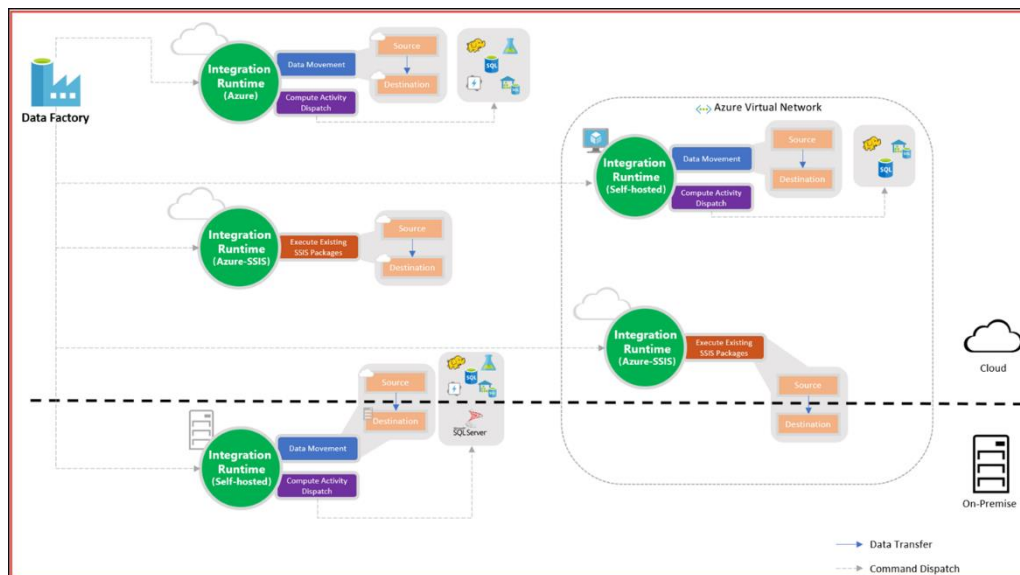
Data movement

Activity dispatch

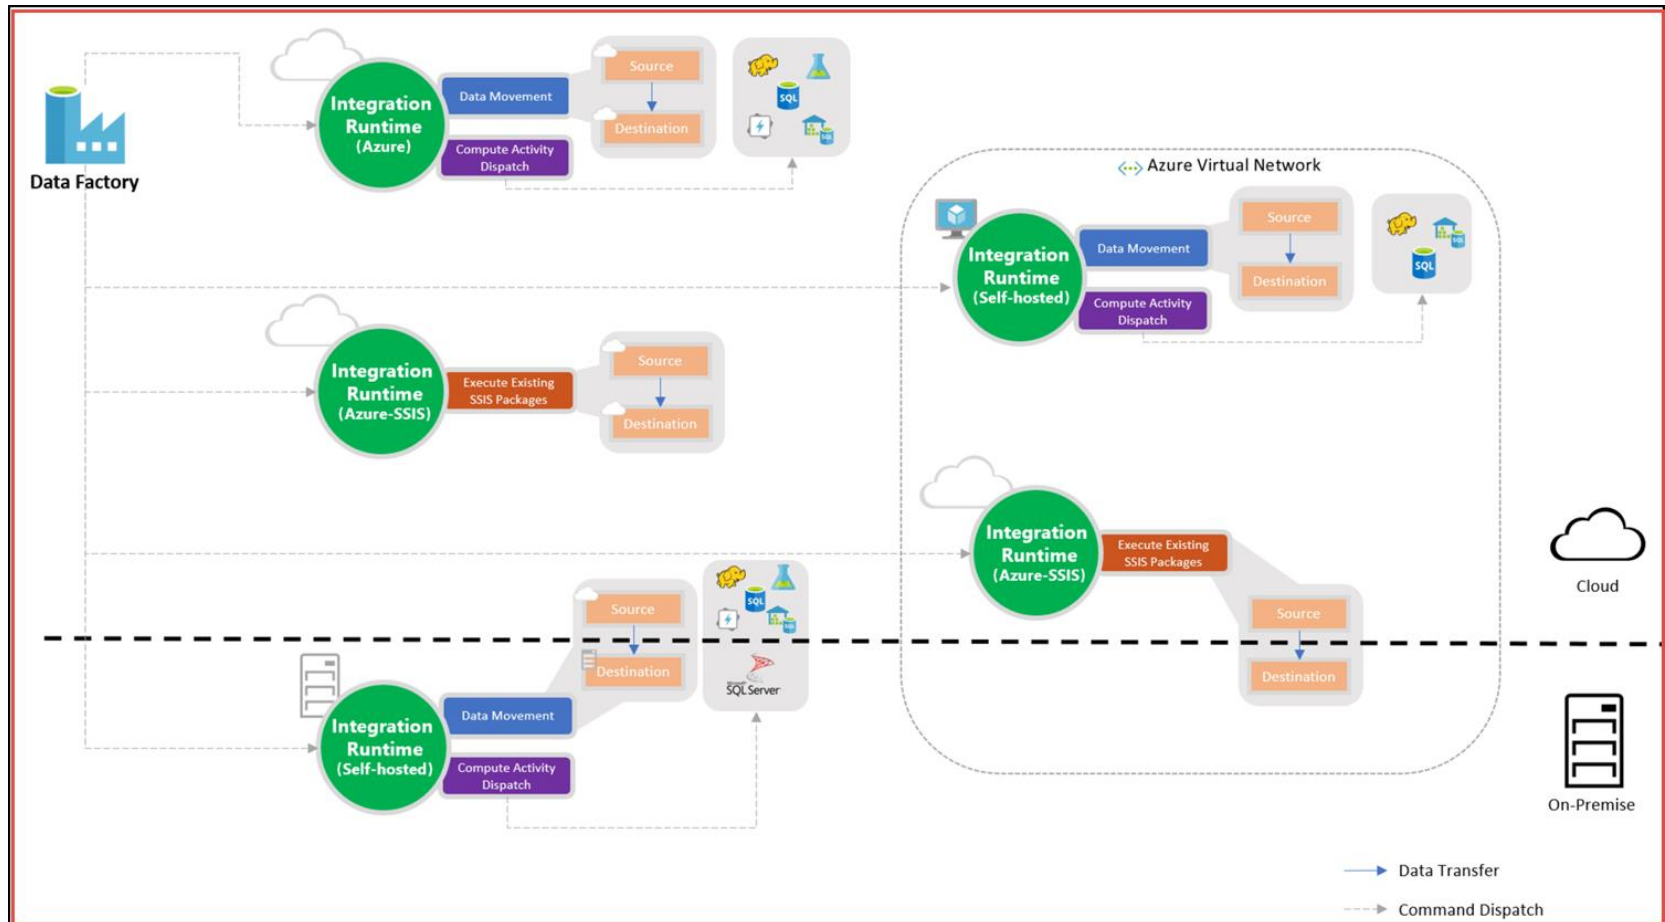
Azure-SSIS

SSIS package execution

SSIS package execution



ADF – Integration Runtimes



ADF – Integration med GIT

- Hvor er min Save knap?
- Vi har kun Publish hvis vi ikke har et GIT repository

